


Connector + diffusion training diagnostics — experimental


Recorded optimizer steps 1–100 (100 logged) | Cumulative examples consumed: 100 / 9,647 training records

Connector and full diffusion transformer are trained. PRISM, VAE and native conditioner stay frozen.


A Optimization loss — stochastic timestep and noise



B Fixed-noise probes — matched captions, unchanged cohorts



C Caption-conditioning control — identical target, timestep and noise



At the last train probe, 16/16 examples had reached the optimizer. Validation examples never enter optimization.

Full-validation markers show the complete split. Fixed curves retain their original IDs/seeds.

Positive caption gaps favor matched captions; they do not establish image quality. Panels A and B use separate y-axis scales.