


PRISM connector — balanced region alignment

1,000 optimizer steps | 32 train / 8 held-out captions | 4,000 exposures (125 passes over the training subset)

Only the connector is trained. PRISM, the native teacher, and the original generator remain frozen.



Features are compared after the frozen DiT caption RMSNorm; fixed-cohort means weight each caption equally.

Objective: one third each for prefix, content, and suffix MSE, independent of region token counts. Fresh optimizer from the strict v1 aligned connector.

Feature matching is not evidence of image-generation quality; paired flow controls and sampled images are separate checks.