


PRISM connector feature alignment

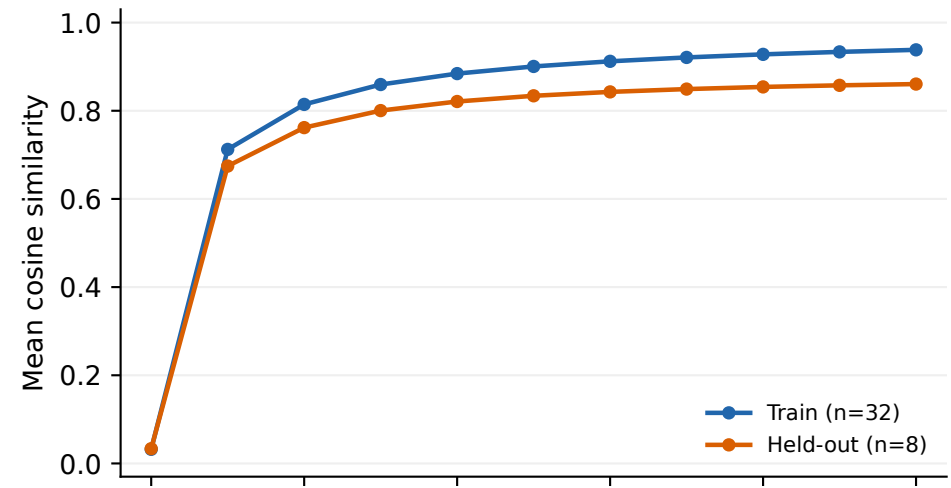
1,000 optimizer steps | 32 train / 8 held-out captions | 4,000 exposures (125 passes over the training subset)

Only the connector is trained. PRISM, the native teacher, and the original generator remain frozen.

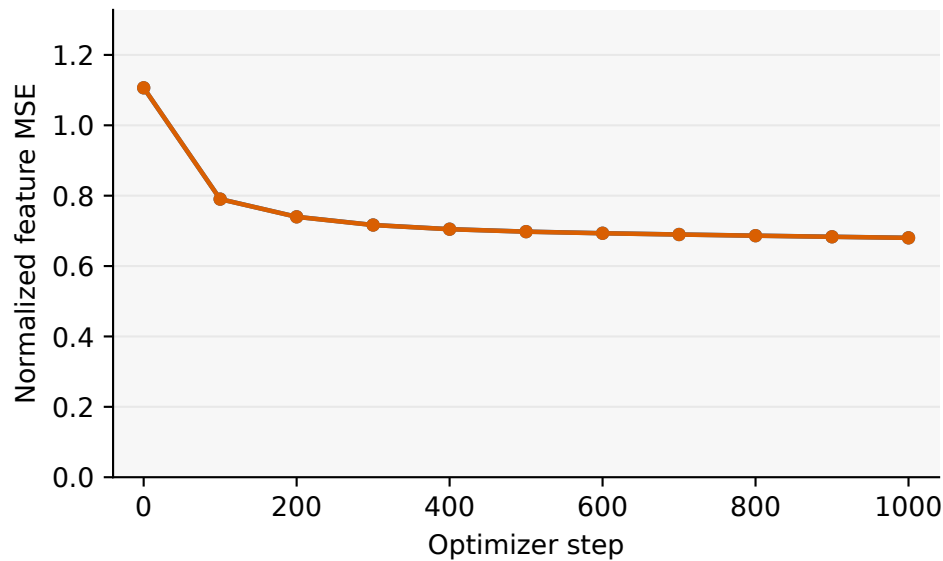
A Caption-content MSE — training objective



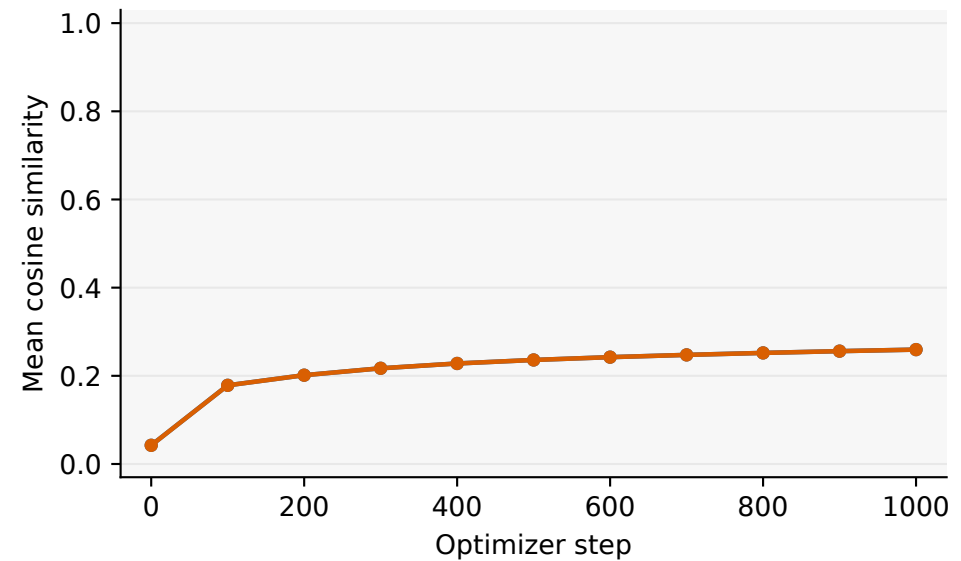
B Caption-content cosine similarity



C Template MSE — excluded from loss



D Template cosine — excluded from loss



Features are compared after the frozen DiT caption RMSNorm; fixed-cohort means weight each caption equally.

Template positions are excluded from optimization; their feature similarity is diagnostic only.

Feature matching is not evidence of image-generation quality; paired flow controls and sampled images are separate checks.